

## 1 Transformers

### Token Embeddings

$\mathbf{I} \in \mathbb{R}^{\text{voc} \times N}$ : Input Sequence -  $\mathbf{M}_{\text{emb}} \in \mathbb{R}^{\text{dim\_emb} \times \text{voc}}$ : Embedding Matrix -  $\mathbf{M}_{\text{pos}} \in \mathbb{R}^{\text{dim\_emb} \times N}$ : Positional Embedding -  $\mathbf{E} \in \mathbb{R}^{\text{dim\_emb} \times N}$ : Token Embeddings

$$\mathbf{E} = \mathbf{M}_{\text{emb}} \times \mathbf{I} + \mathbf{M}_{\text{pos}}$$

### Self Attention

$\mathbf{W} = [\mathbf{W}_q \ \mathbf{W}_k \ \mathbf{W}_v]^\top \in \mathbb{R}^{3\text{dim\_emb} \times \text{dim\_emb}}$  -  $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{\text{dim\_emb} \times N}$ : Query, Key, Value -  $\mathbf{A} \in \mathbb{R}^{N \times N}$ : Attention Matrix

$$[\mathbf{Q} \ \mathbf{K} \ \mathbf{V}]^\top = \mathbf{W} \times \mathbf{E}$$

$$\mathbf{A} = \text{softmax} \left( \frac{\mathbf{Q}^\top \mathbf{K}}{\sqrt{d_k}} \right) \odot \mathbf{M}_{\text{mask}}$$

$$\mathbf{E}_{\text{att}} = \mathbf{V} \times \mathbf{A}$$

With  $d_k$  being the dimension of the key vectors. E.g. pos\_emb for single head attention

### MLP

$\mathbf{M}_{\text{up}} \in \mathbb{R}^{4\text{dim\_emb} \times \text{dim\_emb}}$ : Up projection -  $\mathbf{M}_{\text{down}} \in \mathbb{R}^{\text{dim\_emb} \times 4\text{dim\_emb}}$ : Down projection

$$E_{\text{mlp}} = \sigma(\mathbf{M}_{\text{down}} \times \sigma(\mathbf{M}_{\text{up}} \times \mathbf{E}_{\text{att}}))$$

With  $\sigma$  being the activation applied elementwise

**Computation:** FLOPS  $\approx 6N \cdot D$  - With  $D$  being the number of training tokens.

## 2 Tokenizers

## 3 Training

### 3.1 Datasets

**Common Crawl:** Database of scraped websites. **WebText:** OpenAI internal dataset. Scraped links from Reddit which received at least 3 karma. Page de-duplication and light-cleaning. 8M Documents, 40GB of text.

**OpenWebText:** Open replication of **WebText**

**C4:** Colossal Clean Crawled Corpus. Filtered version of **Common Crawl**. Discard pages with fewer than 5 sentences and lines with fewer than 3 words. Filter for unwanted keywords. Remove lines with the word Javascript. Remove pages with "lorem ipsum". Remove pages with "{". Deduplicate any three-sentence span occurring multiple times. Filter pages that are not in English.

**The Stack:** Coding dataset.

**PeS2o:** STEM papers.

**DOLMA:** Combination of **Common Crawl**, **C4**, **The Stack**, **Reddit**, **PeS2o**, **Project Gutenberg**, **Wikipedia/Wikibooks**

**The Pile:** 800GB Dataset of Diverse Text. Academic, Internet, Prose, Dialogue and Misc (GitHub, Math ...)

**Dataset Cleaning Pipeline:** Language Filtering  $\rightarrow$  Deduplication (by URL)  $\rightarrow$  Quality Filters  $\rightarrow$  Content Filters  $\rightarrow$  Deduplication (by text overlap)

### 3.2 Evaluation

$$\text{Perplexity: } \exp \left( -\frac{1}{N} \sum_{i=1}^N \log P(t_i) \right)$$

With  $t_i$  being the  $i$ th token in the expected output sequence

### Benchmarks

**Paloma:** Perplexity over a diverse set to text.

**HellaSwag:** QA Benchmark. Most likely answer by perplexity is chosen.

**MMLU:** QA Benchmark. Answers are part of the Prompt. Model can answer A, B, C or D. Most likely token is chosen.

### 3.3 Fine-Tuning

Possible with around 1000 high quality prompts and responses or more.

**RLHF** - Reinforcement Learning from Human Feedback

$$y_1, y_2 \propto \pi_{\text{SFT}}(y \mid x) \quad y_w > y_\ell \mid x$$

$$p^*(y_1 > y_2 \mid x) = \frac{1}{1 + \exp(r^*(x \mid y_2) - r^*(x \mid y_1))}$$

$$\mathcal{L}(r) = -\mathbb{E}[\log \sigma(r(x \mid y_\ell) - r(x \mid y_w))]$$

Where  $r$  is the reward model and  $\mathcal{L}$  is the loss of the reward model.

$$\max_{\pi} \mathbb{E}[r(x, y) - \beta D_{\text{KL}}(\pi(y \mid x) \mid \pi_{\text{ref}}(y \mid x))]$$

$D_{\text{KL}}$  to reduce the deviation from the base model (SFT).

**DPO** - Direct Preference Optimization

$$\mathcal{L}_{\text{DPO}} = -\mathbb{E} \left[ \log \sigma \left( \beta \log \frac{\pi_{\theta}(y_w \mid x)}{\pi_{\text{ref}}(y_w \mid x)} - \beta \log \frac{\pi_{\theta}(y_\ell \mid x)}{\pi_{\text{ref}}(y_\ell \mid x)} \right) \right]$$

## 4 Scaling Laws

$$L(D) \approx \frac{A}{D^{\alpha_D}}$$

$$L(N) \approx \frac{B}{N^{\alpha_N}}$$

$$R(f_{N,D}) = R(f^*) + (R(f_N) - R(f^*)) + (R(f_{N,D}) - R(f_N))$$

$$R(f_{N,D}) = E + L(N) + L(D)$$

Let  $R(f^*)$  be the irreducible error,  $R(f_N) - R(f^*)$  the approximation error of a  $N$ -parameter model and  $R(f_{N,D}) - R(f_N)$  the statistical error  
Parameter Values depend on data and precise model architecture.

#### 4.1 Compute Optimality

$\arg \min L(N, D)$  for a constant compute budget  $C(N, D) = H$  by choosing the best  $N$  and  $D$ .

**Training curve envelope** - Plot training curves ( $x$ :  $\log(\text{FLOPS})$ ,  $y$ :  $L(N, D)$ ), Find the envelope (Minimal loss per FLOP), Plot envelope points twice ( $x$ :  $\log(\text{FLOPS})$ ,  $y$ :  $\log(D)$ ,  $\log(N)$ ), Fit a line to the points and extrapolate to the desired FLOPS to find the optimal  $N$  and  $D$ .

**IsoFLOP Curves** - Train various model sizes with  $D$  such that the final FLOPs are constant, Repeat for different final FLOPs, Plot final loss ( $x$ :  $\log(N)$ ,  $y$ :  $L(N, D)$ ), Locate optimal model size for a given compute budget (loss valley), Plot optimal models ( $x$ :  $\log(\text{FLOPS})$ ,  $y$ :  $\log(D)$ ,  $\log(N)$ ) and extrapolate.

**Parametric fit** - Fit parametric Risk function  $R(f_{N,D})$  to the training results (Training like IsoFLOP), Plot contours ( $x$ :  $\log(\text{FLOPS})$ ,  $y$ :  $\log(N)$ ), Fit line such that it goes through each iso-loss contour at the point with the fewest FLOPs. Extrapolate to desired FLOPs. Alternative - Plot isoFLOP slice ( $x$ :  $\log(N)$ ,  $y$ :  $L(N, D)$ ), plot parametric risk for desired compute budget, Locate the minimum.

**Key finding:** For compute optimal training  $D$  should scale proportionally with  $N$  - Compute Optimality doesn't consider inference cost

## 5 Ensembles and MoE

### 5.1 Ensembles

**Linear interpolation** -  $P(y | x) = \sum_m P_m(y | x)P(m | x)$

**Log-linear interpolation** -  $\text{softmax}(\sum_m \log P(y | x)\lambda_m(x))$

With  $P_m(y \mid x)$  being the output of the model  $m$  and  $P(m \mid x)$  the "reliability" of the model given the Input.  $\lambda_m$  is an interpolation coefficients for the model  $m$ . **Parameter Averaging** - Calculate model weights by accumulating (average, weighted average ...) weights from multiple models

5.2 MoE - Mixture of Experts

**Gaussian mixture model** -  $p(y \mid x) = \sum_k p_{\theta_k}(y)p_k(x)$

Where  $p_{\theta_k}$  is a gaussian distribution parametrized by  $\theta_k$  giving the propability of the output  $y$ .  $p_k$  is the probability of that distribution given the input  $x$ .

Allows the approximation of more complex distributions using only simple distributions (gaussian and logistic)

Routing

Shazeer - Mixtral - Switch routing -

6.1 Multicore Processing

GPUs are optimized for performing the same operation on different data points simultaneously.

**SIMD** - single-instruction multiple-data

**Amdahl's law** - Let the non-parallizable part of a program take a fraction  $s$  of the time, then  $m$  workers can result in a speedup of:

$$\frac{1}{s + \frac{1-s}{m}}$$

Floating Point Numbers

	Sign	Exponential	Mantissa
<b>FP32</b>	1 Bit	8 Bits	23 Bits
<b>FP16</b>	1 Bit	5 Bits	10 Bits
<b>BF16</b>	1 Bit	8 Bits	7 Bits

**BF16** - **FP32** range with **FP16** precision

**Error:**  $|x - \tilde{x}| \leq \epsilon/2 \mid x \mid$  with  $\epsilon \neq 0$

6.1.1 CUDA

thread - core; thread block - streaming multiprocessor (SM); kernel grid - CUDA-capable GPU

**CUDA Thread** - Smallest compute entity

**CUDA Block** - Group of threads (up to 1024)

**Streaming Multiprocessor** - Executes one **CUDA Block**

6.2 Distributed Computing

6 Scalable Computing

7 Vision Models